MLA / GQA · HISTORICAL LOSS CURVES

MLA validation keeps improving through 1,200 updates

Peak LR 0.00125 · online model losses · three fresh seeds per recipe · no smoothing or new training

Width 256 · separate campaigns

GQA seeds 2811–2813; MLA 2911–2913

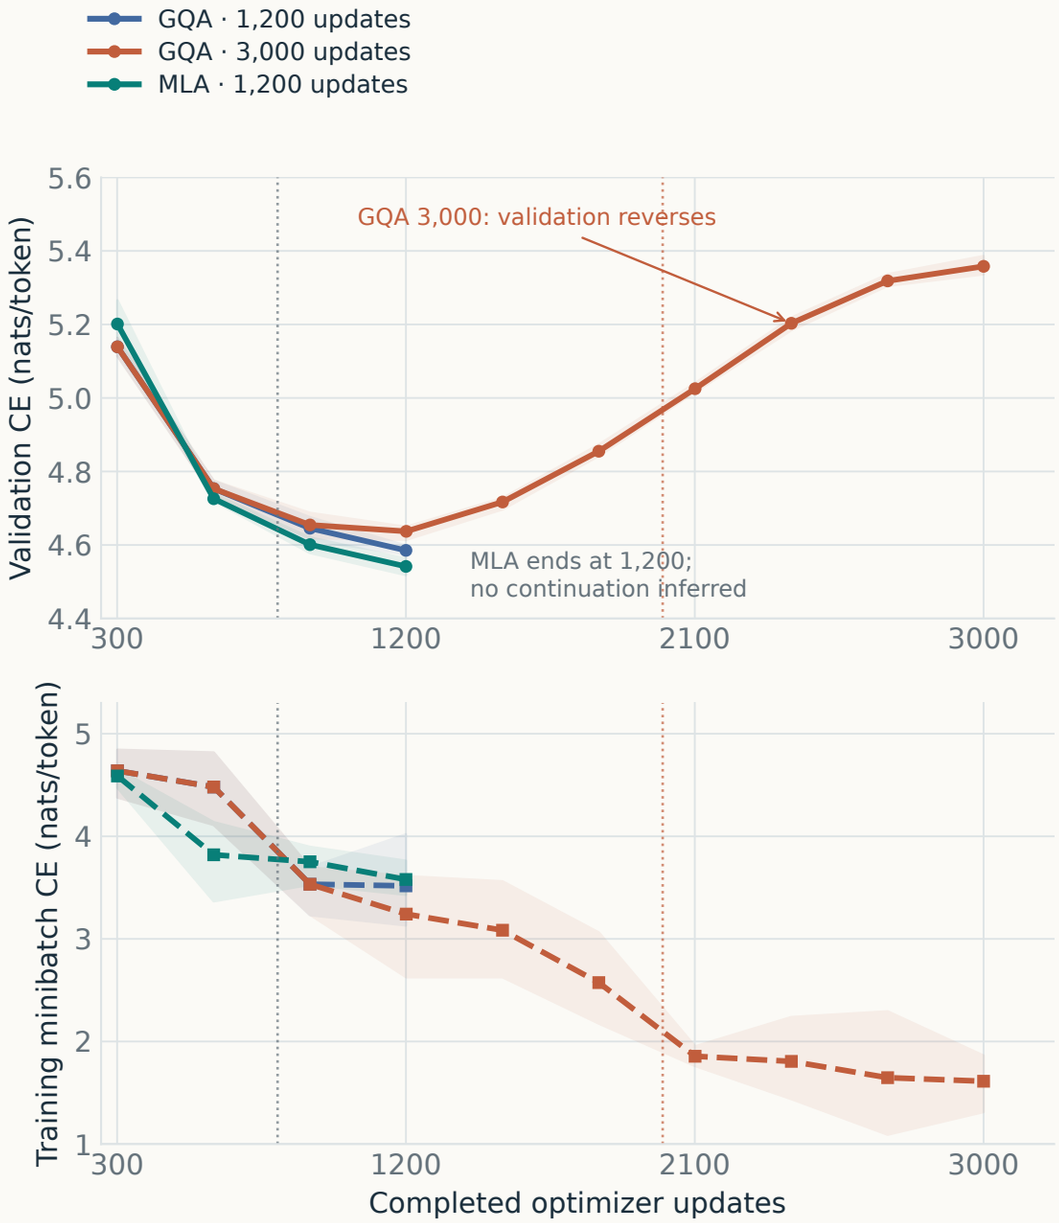

Width 512 · direct paired comparison

Both families: seeds 3411–3413

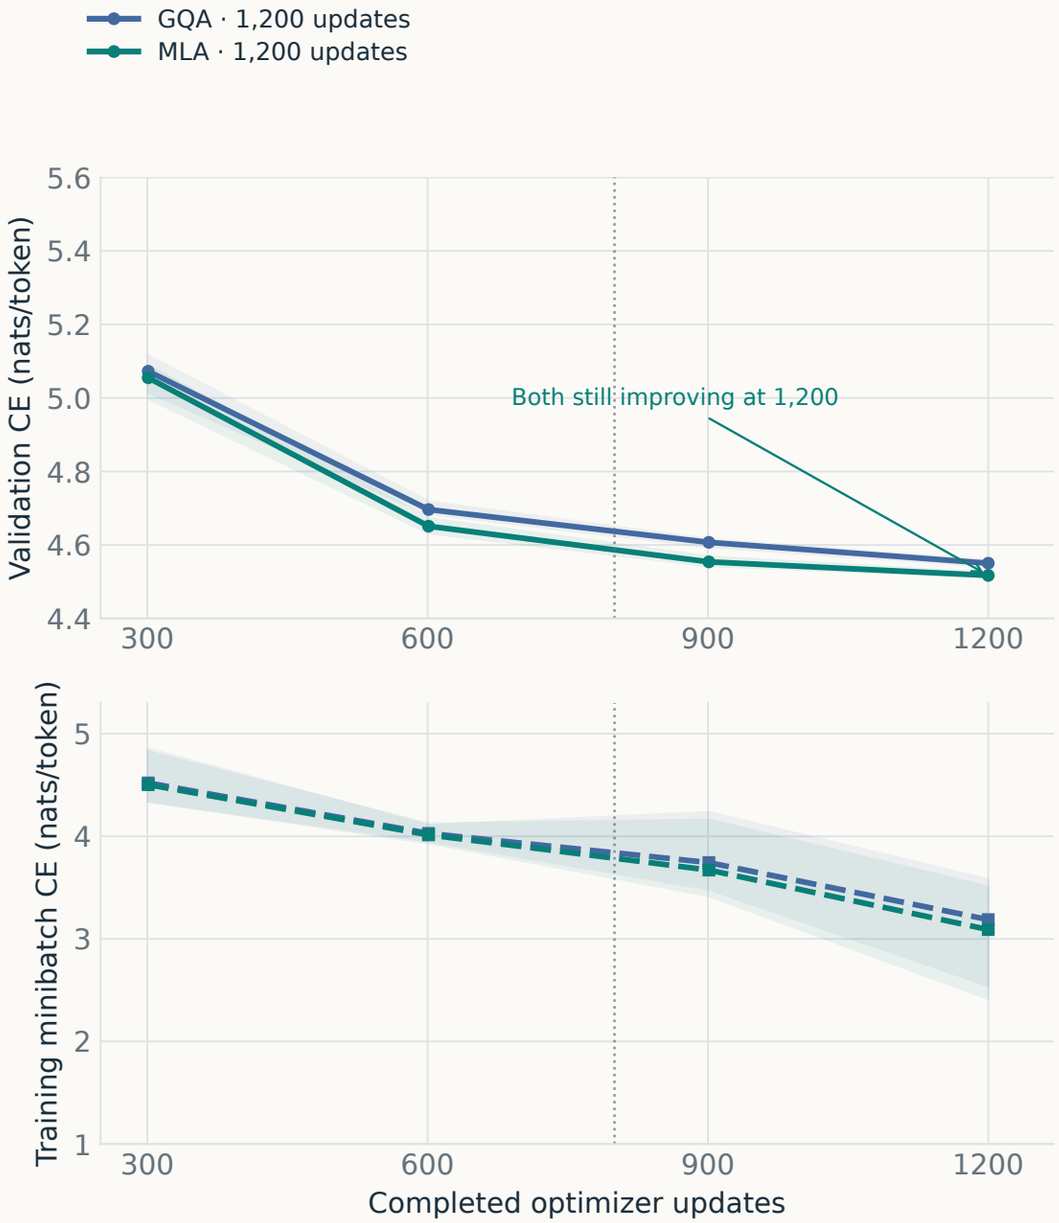

Detail begins at update 301; page 2 includes startup. Shading = range across three seeds, not confidence intervals.
Validation uses the same 10 sampled batches at each log; training uses one noisy minibatch. These are not full-validation EMA curves.
Dotted lines: configured cooldown/EMA start, 800 for short runs and 2,000 for long GQA. D256 is descriptive, with unmatched MLA/GQA seeds.
No MLA 3,000-update curve is available in these comparison campaigns. Improving through 1,200 does not establish resistance to later overfitting.

MLA / GQA · HISTORICAL LOSS CURVES

Complete histories: startup and all individual runs

Peak LR 0.00125 · online model losses · three fresh seeds per recipe · no smoothing or new training

Width 256 · separate campaigns

GQA seeds 2811–2813; MLA 2911–2913

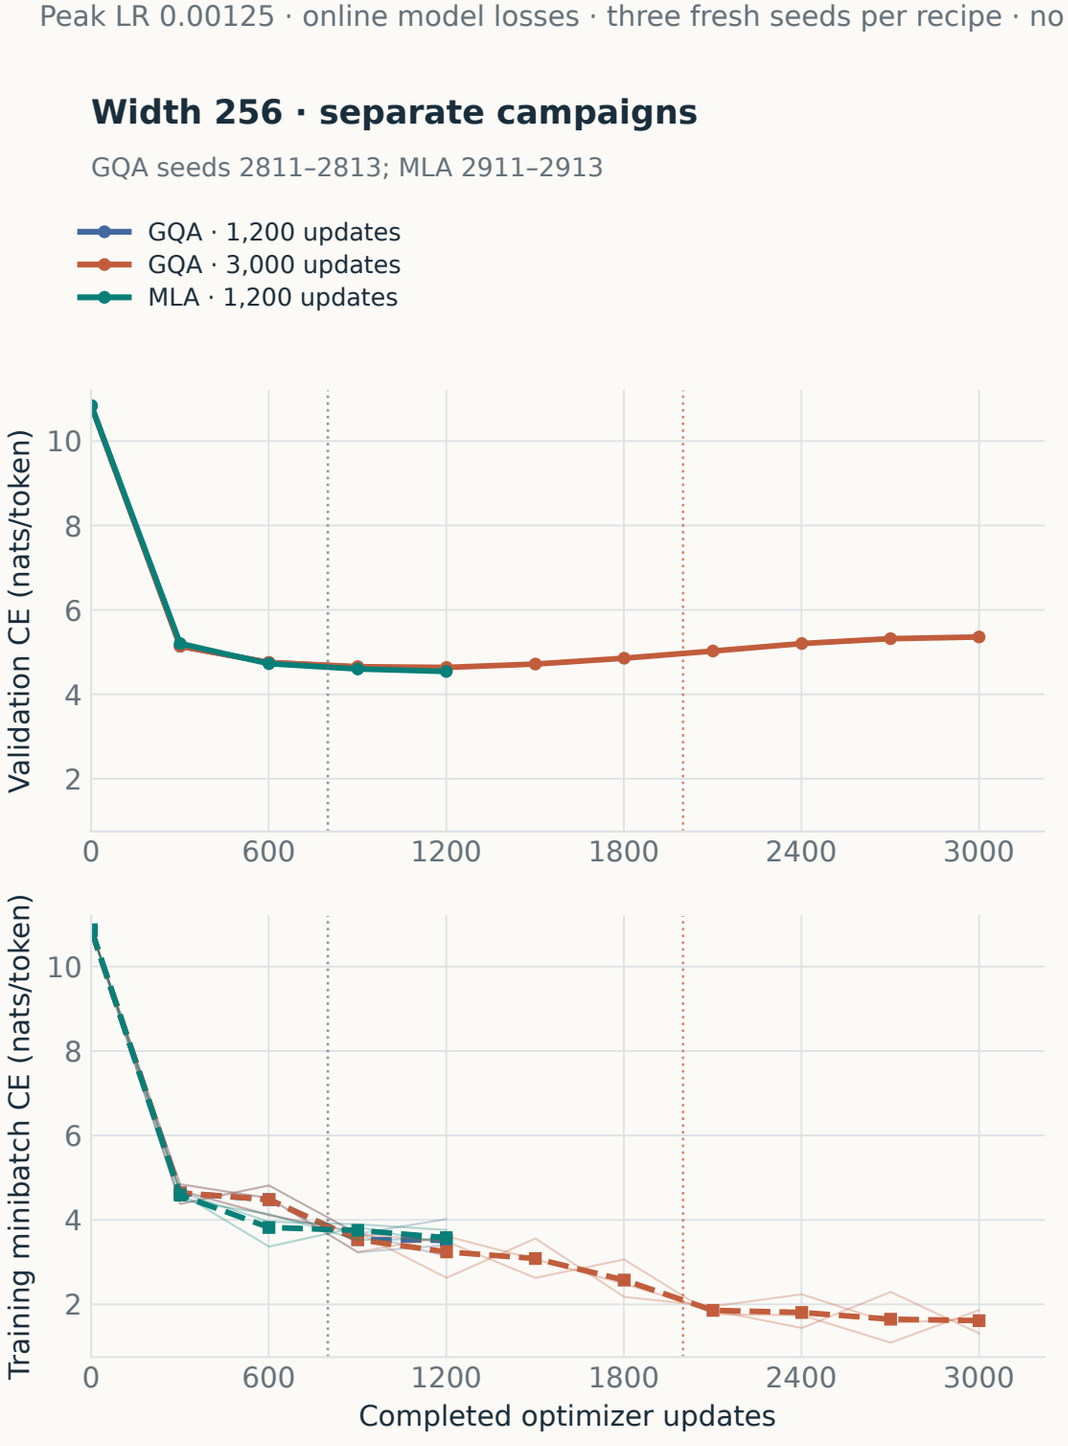

Width 512 · direct paired comparison

Both families: seeds 3411–3413

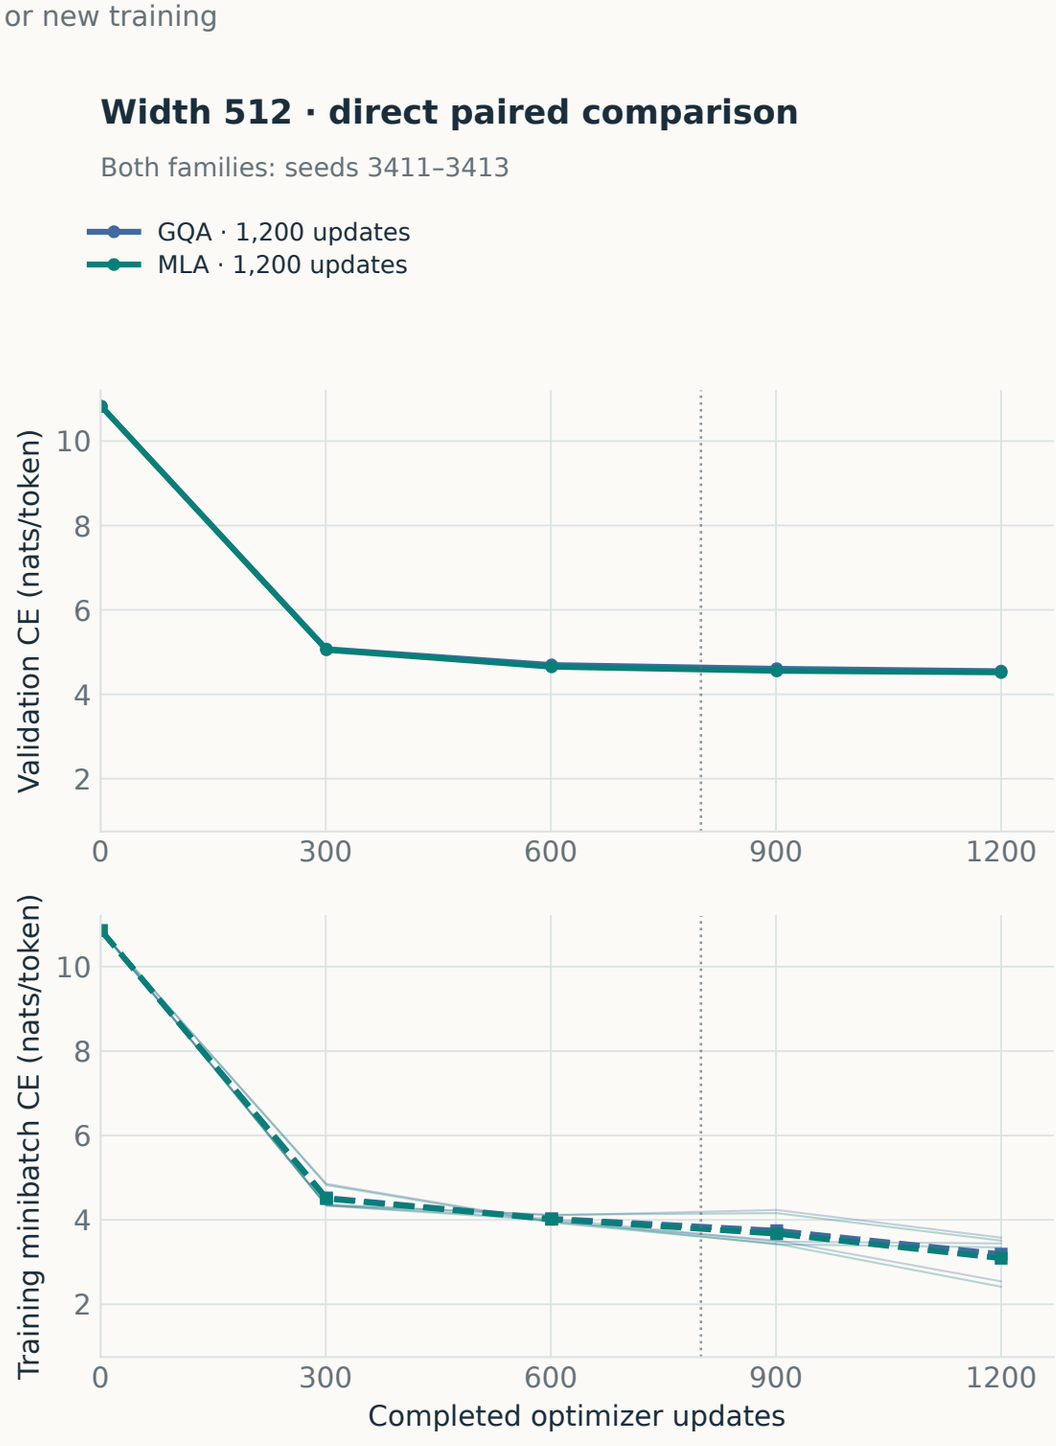

Thin lines = individual seeds; bold lines = three-seed means. Startup CE is about 10.8. Lines connect the actual sparse observations.
Validation uses the same 10 sampled batches at each log; training uses one noisy minibatch. These are not full-validation EMA curves.
Dotted lines: configured cooldown/EMA start, 800 for short runs and 2,000 for long GQA. D256 is descriptive, with unmatched MLA/GQA seeds.
No MLA 3,000-update curve is available in these comparison campaigns. Improving through 1,200 does not establish resistance to later overfitting.